

::ISUG

International Sybase Users Group

techcast series

MDA Tables in ASE—Tips and Tricks

April 18, 2006

SYBASE®

::ISUG

International Sybase Users Group

techcast series

MDA Tables in ASE—Tips and Tricks



Peter Dorfman
Senior Staff Software Engineer
Sybase, Inc.
peter.dorfman@sybase.com


SYBASE®

::ISUG

International Sybase Users Group

techcast series

MDA Tables in ASE—Tips and Tricks



Rob Verschoor
Senior Evangelist
Sybase, Inc
robv@sybase.com

SYBASE®

- Quick introduction to MDA tables
- Possible applications of MDA tables
 - What's that application doing?
 - Identifying unused indexes
 - Identifying 'hot' tables
- Historical MDA tables
- Archiving historical MDA table data
- Performance impact of MDA tables
- Counter wrap
- Analyzing stored procedure activity
- Miscellaneous topics
- Recent Enhancements
- Q&A

Quick Introduction to MDA Tables

- MDA tables were introduced in ASE 12.5.0.3
 - MDA = Monitoring and Diagnostic Access
 - also called “monitoring tables”
- 39 proxy tables in ‘master’ database (**35 in ASE 12.5**)
 - monSysSQLText, monObjectActivity, monCachedObject (etc.)
 - Can be accessed with regular SQL statements
 - When queried, tables are created on-the-fly from memory structures
 - No data is stored in **master** database
 - The proxy tables may also be created in a different database than **master**
- Must be installed: ‘installmontables’ script
- No license needed: included in ASE base product
- Only small performance impact on ASE (<5%)

Quick Introduction to MDA Tables

- MDA tables provide access to low-level monitoring data
 - Resource usage in ASE per table/query/entire server
 - Current activity in ASE per spid/query/procedure/table
 - Recent activity: recently completed statements, with the resources they required
- Some examples of practically relevant information:
 - Amount of memory occupied in the data cache by each table or index
 - Most frequently used tables/procedures
 - Top-N queries for CPU, I/O, elapsed time,...
 - Find unused indexes
 - SQL text of currently executing/recently executed statements
 - Automatically kill user processes that have been idle for more than X minutes
 - Provide server status information even when **tempdb** is full

What *ARE* the MDA Tables?

- The MDA tables are proxy tables
- Defined in the master database
- Defined in installmontables script
- Consume no space
- Table definitions map to ASE RPC calls

```
create existing table monProcessLookup (  
    SPID                smallint,  
    KPID                int,  
    Login               varchar(30) NULL,  
    Application         varchar(30) NULL,  
    ClientHost          varchar(30) NULL,  
    ClientIP            varchar(64) NULL,  
    ClientOSPID         varchar(30) NULL,  
)  
external procedure  
at "loopback...$monProcessLookup"
```

Reference to
Local server



The diagram consists of two grey rectangular boxes. The box on the left contains the text 'Reference to Local server'. The box on the right contains the text 'RPC Call'. Two curved arrows originate from these boxes and point towards the 'at' clause in the SQL code block above. The arrow from the 'Reference to Local server' box points to the 'loopback...' part of the string, and the arrow from the 'RPC Call' box points to the '\$monProcessLookup' part of the string.

RPC Call

Quick Introduction to MDA Tables

- For more MDA basics, and a brief discussion of all tables:
 - See presentations from past Techwave conferences (www.sypron.nl/mda)
- In this presentation:
 - We want to go one step further than just the basics
 - Look at practical applications of MDA tables
 - Things that are useful for you as a DBA

Possible Applications of MDA Tables

What's That Application Doing? SQL Text

- Does this sound familiar?
 - A third-party 'black box' application runs on your ASE server
 - You have the feeling it sometime slows down the entire server...
 - ... but you don't know which queries it is sending to ASE
- Classic solutions:
 - Use "cmdtext" auditing to intercept the application's T-SQL commands
 - Use traceflag 11202 (writes all incoming client language to the errorlog)
 - Use third-party tools to find T-SQL commands by intercepting network packets
 - dbcc sqltext()
 - ...but all these methods have significant limitations or drawbacks

What's That Application Doing?

- Solution: MDA tables:
monProcessSQLText & monSysSQLText
- monProcessSQLText: currently executing SQL
- monSysSQLText: recently executed SQL, now completed
 - Historical table
 - Lets you “look back” in time
 - By copying rows regularly into an ‘archive’ table, complete history can be preserved

What's That Application Doing? Statement Statistics

- Also: monSysStatement: info about completed SQL statement
 - Number of logical I/Os
 - Number of physical I/Os
 - Number of network packets sent/received
 - Number of milliseconds of 'waiting time' during statement execution
 - Exact starttime & endtime of execution
 - Not the SQL Text itself; for this, see **monSysSQLText**
 - Historical table
 - Lets you "look back" in time
 - By copying rows regularly into an 'archive' table, complete history can be preserved

Measuring Statement I/O and CPU Time

- Most I/O intensive statement

```
select * into #ts from master..monSysStatement
```

```
select KPID, BatchID, LineNumber, LogicalReads, Elapsed = datediff(ms,  
    StartTime, EndTime) from #ts where LogicalReads > 10000  
    order by 4 desc
```

KPID	BatchID	LineNumber	LogicalReads	Elapsed
574619857	9	13	5509405	68613
575602932	10	3	360656	7956
575209751	10	3	86241	606
575406493	10	3	59546	983
576258422	2	62	39963	223
575275534	1	62	39959	216
575930476	2	62	39955	250
332857800	65454	1	15884	1600
575275534	9	1	12758	176

Capturing SQL Text

```
select * into #tsql from master..monSysSQLText
```

```
select SQLText from #tsql where KPID= 574619857  
order by BatchID, SequenceInBatch
```

SQLText

```
select  admnr = lvd.id_lm_adres,  
        lvd.id_logmiddel  
from  
        logmi.dbo.lm_voorraad lvd,  
        ravar.dbo.adm_relatie adm  
where  
        lvd.cdsys_lm_adrestype = "A"  
and     lvd.cdsys_lm_opslagstat = "O"  
and     lvd.id_lm_adres        = adm.admnr  
and     adm.dat_ingang         <= @vandaag  
and     (adm.dat_einde         >= @vandaag  
or      adm.dat_einde         = null)
```

Monitoring Table Activity

- Monitoring activity on a specific table

```
select SQLText from master..monSysSQLText  
where SQLText like '%MyTable%'
```

- Also handy for RepServer DBAs:
 - Quick way to figure out exactly which SQL is executed against your replicate DB
 - Especially handy when developing/debugging custom function strings

Monitoring Index Utilization

- Have you ever wanted to see
 - Which indexes are never used?
 - How frequently they are used?
 - How many inserts, deletes, updates, physical or logical I/O they incur?
- monOpenObjectActivity table provides:
 - Table usage count
 - Index usage count
 - Last used dates
 - Physical, logical I/O
 - Row-level insert/delete/update counts
 - Lock wait counts for tables and indexes
- NOTE: Statistics are reset when server is booted or object descriptor is reused in memory.

Monitoring Index Utilization

monOpenObjectActivity

jisql oban:5003

Go Server: oban:5003 Database: master

Input window

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID), IndID = IndexID, UsedCount, LastUsedDate, OptSelectCount, La:
from monOpenObjectActivity
order by UsedCount
```

Output window using table

Database	Table	IndID	UsedCount	LastUsedDate	OptSelectCount	LastOptSelectDate
master	spt_monitor	0	7	2004-01-24 16:05:05.15	7	2004-01-24 16:05:05
pubs2	titleauthor	0	8	2004-01-25 09:33:15.063	8	2004-01-25 09:33:15
tempdb	#1_____00000...	0	10	2004-01-24 15:46:57.156	8	2004-01-24 15:46:57
master	monCachedObject	0	13	2004-01-24 16:40:38.163	13	2004-01-24 16:40:38
master	monOpenObjectActivity	0	13	2004-01-25 09:55:41.063	13	2004-01-25 09:55:41
master	monSysStatement	0	16	2004-01-24 17:26:48.163	3	2004-01-24 15:33:07
master	spt_values	1	22	2004-01-24 15:31:03.07	23	2004-01-24 15:31:03
master	monProcessSQLText	0	36	2004-01-24 16:12:40.153	11	2004-01-24 16:03:14
pubs2	authors	1	36	2004-01-25 09:31:41.663	36	2004-01-25 09:31:41
pubs2	salesdetail	3	40	2004-01-25 09:55:37.063	40	2004-01-25 09:55:37.063
pubs2	sales	0	41	2004-01-25 09:51:40.066	41	2004-01-25 09:51:40.066
pubs2	titles	1	50	2004-01-25 09:53:47.066	50	2004-01-25 09:53:47.066

Status

Table and Index Usage

- Counts
- Dates


```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID),
      IndID = IndexID, UsedCount, LastUsedDate, OptSelectCount, LastOptSelectDate
from master..monOpenObjectActivity
order by UsedCount
```

Monitoring Index Utilization: Unused Indexes


Find Indexes in a Database that have not been used since the server was started

```
select DB = convert(char(20), db_name()),  
       TableName = convert(char(20), object_name(i.id, db_id())),  
       IndexName = convert(char(20), i.name),  
       IndID = i.indid  
from master..monOpenObjectActivity a,  
     sysindexes i  
where a.ObjectID =* i.id  
     and a.IndexID =* i.indid  
     and (a.UsedCount = 0 or a.UsedCount is NULL)  
     and i.indid > 0  
     and object_name(i.id, db_id()) not like "sys%"  
order by 2, 4 asc
```

Outer join finds all indexes



Report indexes that
Have not been used



Monitoring Table Usage

monOpenObjectActivity

jisql oban:5003

Go Server: oban:5003 Database: master

Input window

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID), IndexID, RowsInserted, RowsDeleted, RowsUpdated, LockWaits
from monOpenObjectActivity
order by RowsInserted desc
```

Output window using table

Database	Table	IndexID	RowsInserted	RowsDeleted	RowsUpdated	LockWaits
testdb	t1	0	63	51	11	1
tempdb	#16	00000120016207904	0	13	0	0
tempdb	#11	00000120016207904	0	12	0	0
tempdb	#12	00000120016207904	0	12	0	0
tempdb	#14	00000120016207904	0	12	0	0
tempdb	#15	00000120016207904	0	12	0	0
master	syblicenseslog	0	3	0	0	0
pubs2	sales	0	0	0	0	0
pubs2	titles	0	0	0	0	0

Status

Per Table

- Inserts
- Deletes
- Updates
- Lock Waits

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID),
      IndexID, RowsInserted, RowsDeleted, RowsUpdated, LockWaits
from monOpenObjectActivity
order by RowsInserted desc
```

Identifying 'Hot' Tables

- What makes a table "hot"?
 - Logical reads?
 - Physical reads?
 - Number of queries?
 - Lock usage?
- monOpenObjectActivity reports a number of measures of table and index activity
- Example

```
select * into #t
from master..monOpenObjectActivity
go
```

```
select TableName = object_name(ObjectID, DBID), IndexID,
       LogicalReads, PhysicalReads, Operations, LockWaits
from #t
order by 3 desc
go
```


Identifying 'Hot' Tables

TableName	IndexID	LogReads	PhysReads	Operations	LockWaits
products_tb	0	282294	9043	609	97
products_tb	2	36450	0	0	0
cust_tab	0	12315	0	17	2
cust_tab	2	239	0	0	0

Understanding and Using Historical Tables

Using Historical Tables

- Which MDA tables are “historical” tables?
- What are Historical Tables?
- How do they work?
- What is the correct size to configure them?
- Archiving historical table data
- Tips on using historical tables



Which Tables are Historical Tables?

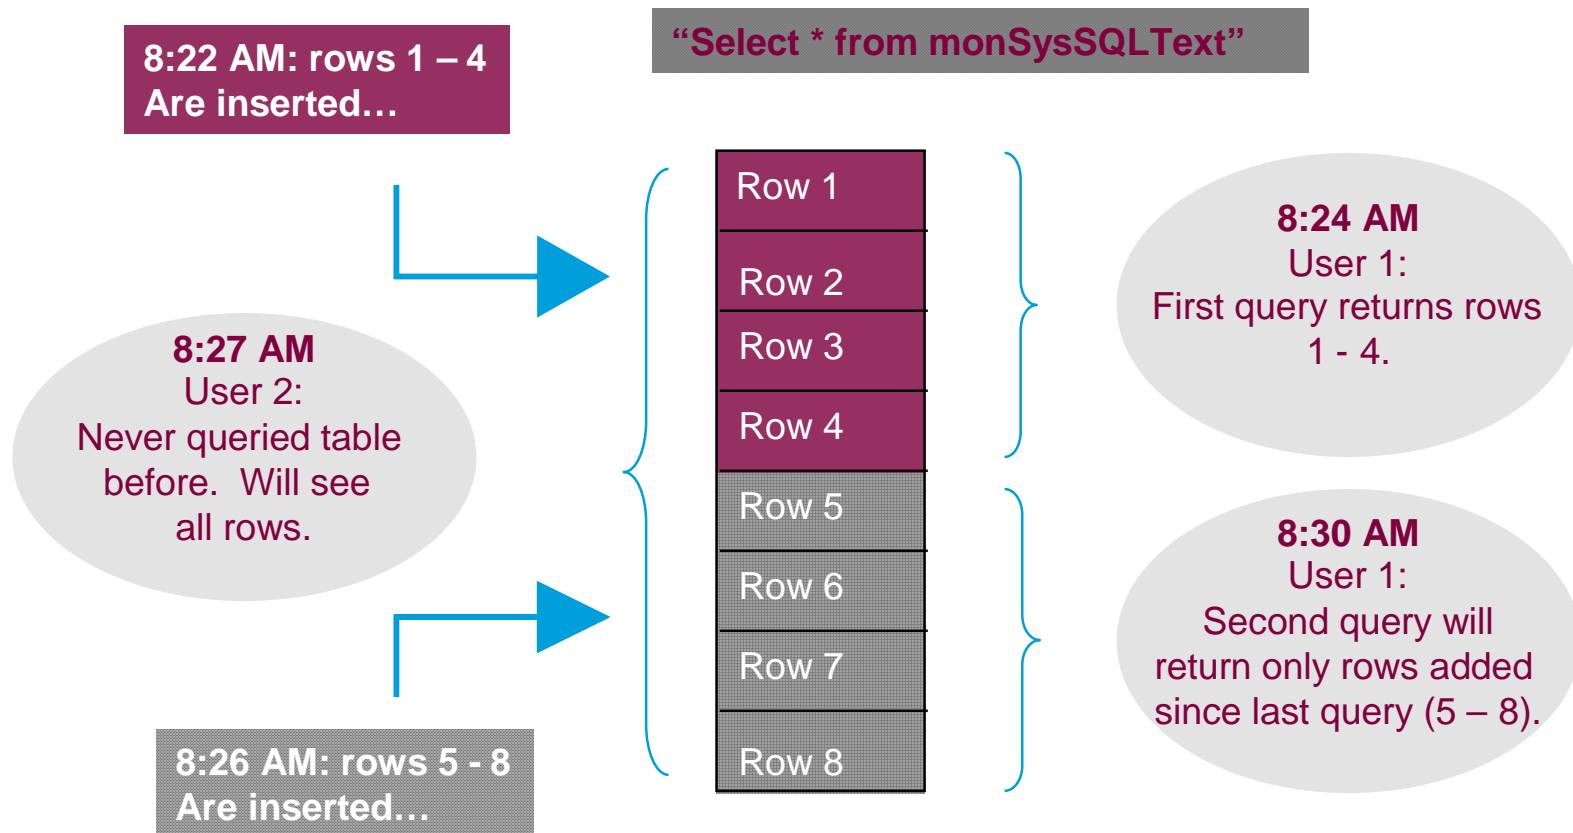
- monSysSQLText
 - Records every SQL command executed on the server
- monSysPlanText
 - Records the Query Plan for every SQL command executed on the server
- monSysStatement
 - Reports the statistics for every statement within every query, batch, stored procedure, trigger, etc. executed on the server
- monErrorLog
 - Records every row written to the server errorlog
- monDeadLock
 - Records information on every deadlock that occurs on the server

What are Historical Tables?

- The historical MDA tables contain a record of “events” within the ASE
 - E.g., SQL submitted for a query, a statement executed within a batch, error message added to the errorlog
- The data for these tables is stored in memory in fixed-sized arrays
 - Size is configurable using `sp_configure`
- Data in Historical tables is transient
 - The arrays are managed as “ring buffers”: After the last entry in the array is written the first entry will be overwritten
- Historical tables are “stateful.”
 - The ASE remembers which records a process has already seen
 - Subsequent queries on same table will return only new records
- Why are they stateful?
 - This allows applications to accurately collect or “drain” the rows in these tables without finding duplicates

Queries on Historical Tables

- The ASE maintains the connection's *currency* in the MDA table
- Currency is reset for each new connection



Setting the Size of Historical Tables

- Depends on which table you are configuring
- These sp_configure parameters determine the number of rows in the historical tables
 - errorlog pipe max messages
 - plan text pipe max messages
 - sql text pipe max messages
 - statement pipe max messages
 - deadlock pipe max messages
- The value of the parameter is the number of rows *per engine*
- Correct size depends on
 - Rate at which rows are written to table
 - Frequency with which queries will be run against the table
- For example:
 - 2 engines
 - 5000 rows per minute per engine
 - Select * from monSysStatement every 5 minutes
 - Statement pipe max messages should be greater than or equal to 25000
 - Result set size??? (50000 rows!)
- Errorlog and deadlock pipes are usually much smaller than plan text, sql text and statement pipes

Rate x Frequency = Size
E.g.: 5000/min x 5 min = 25000

Reasonable size on
busy system??
Could be >> 100000

Tips on Using Historical Tables

- Do not use in subqueries or joins
- Save contents of tables to an archive table or database for analysis
- When collecting long-term data, archive data on a regular basis and size tables to avoid data loss
- How do you know whether the table for the buffer has wrapped?
 - If # of rows returned = size of buffer * # of engines
 - In other words, if you get the entire size of the buffer, some rows were probably lost
 - Only a “rule of thumb”
 - Currently, it is not possible to determine how many rows were lost

Archiving Historical Table Data

Archiving Historical Table Data


- Why is an historical archive useful?
- Because data in historical tables is transient
 - Capture data for later or detailed analysis
- Because repeated queries on historical tables will not return the same rows
 - Even in subqueries or joins
 - Makes analytical queries directly against MDA tables difficult
- Solution
- Data from historical tables should be moved to separate storage for analysis
- Create a monitoring data repository for historical diagnostics or capacity planning

Archiving Historical Table Data

- A possible approach: a 'collector' stored proc which frequently extracts data from the MDA tables
 - 'sp_mda_collect'
 - sp_mda_collect 'start' [, 'hh:mm:ss'] -- runs in a loop (default interval = 30 sec.)
 - sp_mda_collect 'stop' -- run from a different session, stops the original procedure
 - sp_mda_collect 'status' -- displays #rows saved in archive tables
 - (download from www.sypron.nl/mda)
 - Uses a separate database to collect the historical data in permanent tables
 - The permanent tables have the same layout as the historical MDA tables
 - Added a composite unique index **with ignore_dup_key** on key columns (SPID, KPID, etc.) to filter out duplicates (in case the proc needs to be restarted...)

Enterprise Monitoring Repository and Center

- To access MDA tables from a remote server
 - Create the MDA proxy tables on a central server
 - Map MDA proxy tables to each monitored server
- Reduces load on monitored ASE servers
- Provides central source of monitoring data for your enterprise
- Allows easy archiving of enterprise data to permanent storage in database on repository server



Creating an Enterprise Monitoring Center

- Create monitoring database on central server
- Copy and edit installmontables script
 - Two options:
 - Create separate monitoring database for each monitored server
 - Add server name to MDA table names to create unique table names for each server within a single database
- Set the use database command to use the correct database
- Use sp_addserver to register the remote monitored server with the repository server
 - Use sp_addexternlogin to coordinate login credentials if needed
- Change the “loopback” server name to the remote server name of the monitored server in your central server
- Copy object ID's and names from remote servers on a regular basis

Modifying installmontables Script

- Creating MDA proxy tables in a separate database for each monitored server

```
...  
...  
use monitor_svrtest1  
go  
...  
...  
  
create existing table  
monProcedureCache (  
           Requests      int,  
           Loads         int,  
           Writes        int,  
           Stalls        int,  
)  
external procedure  
at "svrtest1...$monProcedureCache"  
go
```

Use a separate database for each
Monitored server

Proxy table points to monitored
server

Modifying installmontables Script

- Creating MDA proxy tables in a single database for all monitored server

```
...  
...  
use monitordb  
go  
...  
...  
  
create existing table  
monProcedureCache_svrtest1 (  
    Requests      int,  
    Loads         int,  
    Writes        int,  
    Stalls        int,  
  
)  
external procedure  
at "svrtest1...$monProcedureCache"  
go
```

Database in which all proxy tables
Will be created

Unique table name constructed by a
Appending server name

Proxy table points to monitored
server

Performance Impact of MDA Tables

Performance Impact of MDA Tables

- Two questions
 - Impact of data collection?
 - Impact of querying MDA tables?
- General performance impact: 5% or less
- Depends on a number of factors
 - Configuration of server (e.g., number of engines, memory size, processor speed)
 - Load on server
 - Configuration of Monitoring parameters
- Different monitoring configuration settings have different performance impacts
- Fully enabling all options will have greatest impact

Performance Impact: Configuration Settings

- **Lowest impact**

- Enable monitoring with no other options

⇒ **Tables enabled**

- monEngine, monDataCache, monProcedureCache, monOpenDatabases, monSysWorkerThread, monNetworkIO, monLocks, monCachePool, monIOQueue, monDeviceIO, monProcessWorkerThread, monProcessNetIO

- **Medium Impact Parameters**

- wait event timing
- plan text pipe active
- sql text pipe active
- errorlog pipe active
- deadlock pipe active

⇒ **Tables enabled**

- monSysWaits, monProcessWaits
- monSysPlanText, monSysSQLText, monErrorLog, monDeadLock

- **Highest Impact Parameters**

- statement pipe active
- statement statistics active
- per object statistics
- statement pipe active

⇒ **Tables enabled**

- monOpenObjectActivity, monProcessObject, monProcessActivity, monSysStatement, monProcessStatement

**** Note:** Actual impact depends on system load patterns and configuration

Performance Improvements in ASE 15.0 ESD#2

- Up to ASE 15.0 ESD#1:

```
create existing table monLocks (  
...  
)  
external procedure  
at "loopback...$monLocks"  
go
```

'external procedure':
When querying the MDA table, ASE
creates a 'backdoor' client connection
back into ASE itself
('loopback' = current ASE server)

- As of ASE 15.0 ESD#2:

```
create existing table monLocks (  
...  
)  
materialized  
at "$monLocks"  
go
```



Why?

- Reduces resource usage
- Increases query speed

New option 'materialized'

Does not require the additional 'backdoor'
client connection anymore

Less overhead, better performance:

- no additional connection management
- saves network I/O

Understanding and Handling Counter Wrap

What is Counter Wrap?

- All MDA counter columns are 32-bit signed integers
 - Maximum value is 2147483647
- When signed integers are incremented above maximum value they become negative
 - $2147483647 + 1 \Rightarrow -2147483646$
- Internal adjustments prevent MDA counter values from becoming negative
 - Therefore counter ranges are from 0 to 2147483647
- When the ASE increments an MDA counter past the maximum value it will return to 0 and start increasing again

Handling Counter Wrap with Delta Values

- If counter has wrapped, add difference between start value and maximum value + 1 to the current value of the counter

```
Select CacheName,  
       CacheSearches =  
       case  
         when e.CacheSearches < s.CacheSeaches  
         then  
           (2147483648 - s.CacheSearches) + e.CacheSearches)  
         else  
           (e.CacheSearches - s.CacheSearches)  
       end  
from #cacheStart s, #cacheEnd e  
where s.CacheID = e.CacheID
```

- Again: As long as change in counter values is <= 2147483647, delta values will be accurate

Which MDA Table Columns Can Wrap?

- Not all MDA columns are likely to wrap
 - Some counter values increment slowly
 - Some numeric columns are not counters
- Columns that can wrap pretty quickly
 - **monDataCache**
 - **CacheSearches**
 - **LogicalReads**
 - **monNetworkIO**
 - **BytesSent**
 - **BytesReceived**
 - **monSysWaits**
 - **Waits**
- Others wrap less quickly
 - **monEngine.ContextSwitches**
 - **monNetworkIO.PacketsSent**

sp_sysmon, subqueries, joins...

Using MDA Tables and sp_sysmon

- Monitor Counters are a set of counters used by sp_sysmon and Monitor Server
- Some MDA table columns are derived from Monitor Counters
- sp_sysmon resets the value of Monitor Counters when it starts
- This can have an impact on applications using MDA tables or Monitor Server
- MDA table columns that come from Monitor Counters are documented.
 - Attributes column = “counter, reset”

Using MDA Tables and sp_sysmon

- The ASE 12.5.3 release introduced changes to sp_sysmon that allow it to run without clearing monitor counters
- Also enhanced so that when multiple applications are using monitor counters the collection of monitor data will not be terminated until all applications are finished
- It is safe to run sp_sysmon when using the MDA tables as long as sp_sysmon is run with the 'noclear' option.

```
> exec sp_sysmon "00:01:00", noclear
```

Subqueries, Joins and Self-Joins

- Rule of thumb: Don't use joins or subqueries when querying the MDA tables
- Why? Because the MDA table data is transient and reflects the ASE's instantaneous state, joins and subqueries may not give the expected result.
 - Sequential queries on same table can give different results
- Because of the currency mechanism, self-joins or subqueries involving one of the historical tables more than once will not work.
 - Currency is reset by first query and the same rows will not be seen by the subquery or inner join table
- Solution: Copy MDA table data to a work table or permanent repository when complex analysis is required.

Analyzing Stored Procedure Performance

Analyzing Stored Procedure Performance

- Historical Server provides stored procedure performance information
- MDA tables do not provide a table with historical stored procedure statistics (yet, we're working on it!)
- The monSysStatement table can be used to report this information

```
select * into #s from master..monSysStatement  
where ProcedureID !=0
```

```
select ProcName = isnull(object_name(ProcedureID,  
    DBID), "UNKNOWN"),  
    DBName = isnull(db_name(DBID), "UNKNOWN"),  
    ElapsedTime = datediff(ms, min(StartTime),  
    max(EndTime))  
from #s  
group by SPID, DBID, ProcedureID, BatchID
```

Stored Procedure Statistics

```

1> select ProcName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
2> DBName = isnull(db_name(DBID), "UNKNOWN"),
3> ElapsedTime = datediff(ms, min(StartTime), max(EndTime))
4> from #s
5> group by SPID, DBID, ProcedureID, BatchID
6> order by 3
7> go

```

ProcName	DBName	ElapsedTime
-----	-----	-----
p_sybbugstatus	engcomdb	1096
sybrev_fetch_revstatus	engcomdb	983
p_sybbugstatus	engcomdb	923
p_sybbugstatus	engcomdb	836
p_sybbugstatus	engcomdb	683
p_sybbugstatus	engcomdb	620
p_sybbugstatus	engcomdb	586
p_sybbugstatus	engcomdb	543
p_sybbugstatus	engcomdb	533
p_sybbugstatus	engcomdb	526
...		
...		
0		

Stored Procedure Performance Averages

Aggregate performance statistics can be derived from the output of the previous query

```
/*
** Build a detail table
*/
select ProcName = isnull(object_name(ProcedureID, DBID),
"UNKNOWN"),
       DBName = isnull(db_name(DBID), "UNKNOWN"),
       ElapsedTime = datediff(ms, min(StartTime), max(EndTime))
into #t1
from master..monSysStatement
group by SPID, DBID, ProcedureID, BatchID
having ProcedureID != 0

/*
** Calculate aggregate values
*/
select ProcName, DBName, "Avg" = avg(ElapsedTime),
       NumExecs = count(*)
from #t1
group by ProcName, DBName
order by 3 desc
```

Stored Procedure Performance Averages

Determine average elapsed time and total executions for stored procedures

```
1> select ProcName, DBName, "AvgElapsed" = avg(ElapsedTime),
2> NumExecs = count(*)
3> from #t1
4> group by ProcName, DBName
5> order by 3 desc
6> go
```

ProcName	Database	AvgElapsed	NumExecs
-----	-----	-----	-----
p_sybugstatus	engcomdb	483	32
sn_temp_filters_qts1	qts_db	330	26
sy_resolution_insert	qts_db	260	44
p_sybugreleasematrix	engcomdb	186	21
create_sn_subscriptions	qts_db	108	9
p_sybugsrelease	engcomdb	91	37
sn_temp_filters_qts2	qts_db	83	2
sn_temp_filters_qts4	qts_db	73	11
sn_get_next_key	qts_db	69	5
create_sn_filters	qts_db	65	5
...			
...			

Identifying Poorly Performing Statements

Identify statements within stored procedures consuming greater than average elapsed time

```
/*
** Build work table
*/
select ProcName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
       DBName = convert(char(15), isnull(db_name(DBID), "UNKNOWN")),
       LineNumber,
       ElapsedTime = datediff(ms, StartTime, EndTime)
into #t1
from master..monSysStatement
where ProcedureID != 0
/*
** Calculate aggregate values and find problematic statements
*/
select ProcName, DBName, LineNumber, "AvgElapsed" = avg(ElapsedTime)
from #t1
group by DBName, ProcName, LineNumber
having avg(ElapsedTime) > (select avg(ElapsedTime) from #t1)
order by 4 desc
```

Statements with > Average CPU Time

...

...

ProcName	DBName	LineNumber	AvgElapsed
-----	-----	-----	-----
row_update	qts_db	614	2160
p_sybbugstatus	engcomdb	60	240
row_update	qts_db	147	98
row_insert	qts_db	308	98
e2_CiMember	qts_db	71	77
p_sybbugstatus	engcomdb	56	76
sy_addl_case_update	qts_db	138	70
p_sybbugstatus	engcomdb	125	69
sybrev_report_newcrs	engcomdb	48	30
log_activity	qts_db	155	18
p_sybbugstatus	engcomdb	29	16
p_sybbugstatus	engcomdb	145	15
log_activity	qts_db	90	14

...

...

Most Frequently Used Stored Procedures

```

1> select * into #t1 from master..monSysStatement
2> go
1> select ProcedureName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
2> "Database" = db_name(DBID),
3> "Execs" = count(*)
4> from #t1
5> where ProcedureID != 0
6> group by DBID, ProcedureID
7> order by 3 desc
8> go

```

ProcedureName	Database	Execs
-----	-----	-----
sp_mltypeset	empdb	8138
p_sybugstatus	engcomdb	888
sp_help_rep_agent	sybsystemprocs	462
p_sybugsrelease	engcomdb	205
sn_get_next_key	qts_db	176
create_sn_filter_criteria	qts_db	162
create_sn_subscriptions	qts_db	136
create_sn_filters	qts_db	120

...

...

Recent Enhancements

Enhancements in 12.5.1 and 12.5.2

- 360 columns in 12.5.0.3 (first version of MDA tables)
- 5 new columns in 12.5.1
- 2 new columns in 12.5.2
 - **monProcessObject.TableSize** - table size in Kb
 - **monProcessActivity.WorkTables** - total number of work tables created by the process
- Fixes:
 - milliseconds fixed in **monSysStatement.StartTime / EndTime**
 - can be used to determine the exact duration of each statement (resolution = 3 milliseconds)

Enhancements in 12.5.3

- New columns in 12.5.3
 - **monProcessActivity.ServerUserID** - Login ID (suid)
 - **monProcessSQLText.ServerUserID** - Login ID (suid)
 - **monSysSQLText.ServerUserID** - Login ID (suid)
 - **monProcessProcedures.LineNumber** – Line number in a stored procedure
- **New columns in 12.5.3 ESD#2:**
 - 4 new columns in **monEngine**: **Yields**, **DiskIOChecks**, **DiskIOPolled**, **DiskIOCompleted**
- Fixes:
 - Various small bugs were fixed

New tables and columns in 15.0

- 4 new tables in ASE 15.0
 - **monOpenPartitionActivity**
 - similar to **monOpenObjectActivity** but from a partitions perspective
 - **monLicense**
 - shows active license keys
 - **monProcedureCacheMemoryUsage**,
monProcedureCacheModuleUsage
 - for engineering usage; no useful info for DBAs
- **Various new columns in** monEngine, monCachedObject, monProcessObject, monOpenObjectActivity

Enhancements in 15.0 ESD#2

- monLocks:
 - Report blocking locks; new columns **BlockedBy** and **BlockingState**
- **New CIS option 'materialized' for MDA proxy tables (see earlier slide)**
- **Improved names of wait events in monWaitEventInfo**
 - Clarified many wait event names
 - Removed most duplicate wait event names

New columns in 12.5.4 and 15.0 ESD#2

- 3 new columns in 12.5.4 and 15.0 ESD#2
 - **monSysStatement.RowsAffected** – like @@rowcount
 - **monSysStatement.ErrorStatus** – like @@error
 - **monProcessStatement.RowsAffected** – like @@rowcount



Your Questions are Welcome

Thanks!

robv@sybase.com

peter.dorfman@sybase.com